

Assignment 1

November 1, 2023

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
```

1 GAUSSIAN

1.1 PART I

Data generation

- (a) Use Gaussian distribution with appropriate parameters and produce a dataset with four classes and 30 samples per class:

the classes must live in the 2D space and be centered on the corners of the unit square (0,0), (0,1), (1,1), (1,0),

all with independent components each with variance 0.3.

- (b) Obtain a 2-class train set $[X, Y]$ by having data on opposite corners sharing the same class with labels +1 and -1.
- (c) Generate a test set $[X_{te}, Y_{te}]$ from the same distribution, starting with 200 samples per class.
- (d) Visualize both sets using a scatter plot on a 2-D plane.
- (e) Repeat (a)-(d) for Laplace distribution.

```
[ ]: # TRAIN

train_samples_per_class = 30

# Set the variance for each class
variance = 0.3

n1a = np.random.normal(loc=[0, 0], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
n1b = np.random.normal(loc=[1, 1], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
n2a = np.random.normal(loc=[0, 1], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
n2b = np.random.normal(loc=[1, 0], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
```

```

labels = np.array([1]*train_samples_per_class*2 +
    ↪ [-1]*train_samples_per_class*2)

X_train = np.vstack([n1a, n1b, n2a, n2b])
Y_train = labels

shuffle_idx = np.random.permutation(X_train.shape[0])
X_train = X_train[shuffle_idx]
Y_train = Y_train[shuffle_idx]

# TEST

test_samples_per_class = 200

n1a = np.random.normal(loc=[0, 0], scale=np.sqrt(variance),
    ↪ size=(test_samples_per_class, 2))
n1b = np.random.normal(loc=[1, 1], scale=np.sqrt(variance),
    ↪ size=(test_samples_per_class, 2))
n2a = np.random.normal(loc=[0, 1], scale=np.sqrt(variance),
    ↪ size=(test_samples_per_class, 2))
n2b = np.random.normal(loc=[1, 0], scale=np.sqrt(variance),
    ↪ size=(test_samples_per_class, 2))

labels = np.array([1]*test_samples_per_class*2 + [-1]*test_samples_per_class*2)

X_test = np.vstack([n1a, n1b, n2a, n2b])
Y_test = labels

shuffle_idx = np.random.permutation(X_test.shape[0])
X_test = X_test[shuffle_idx]
Y_test = Y_test[shuffle_idx]

```

```


[ ]: fig, ax = plt.subplots(1, 2, figsize=(16, 6))
ax[0].scatter(X_train[Y_train == 1][:, 0], X_train[Y_train == 1][:, 1],
    ↪ label='Class 1', marker='o')
ax[0].scatter(X_train[Y_train == -1][:, 0], X_train[Y_train == -1][:, 1],
    ↪ label='Class -1', marker='x')
ax[0].set_title('Training set')
ax[0].set_xlabel('X1')
ax[0].set_ylabel('X2')
ax[0].legend()

```

```

ax[1].scatter(X_test[Y_test == 1][:, 0], X_test[Y_test == 1][:, 1],
             ↪label='Class 1', marker='o')
ax[1].scatter(X_test[Y_test == -1][:, 0], X_test[Y_test == -1][:, 1],
             ↪label='Class -1', marker='x')
ax[1].set_title('Test set')
ax[1].set_xlabel('X1')
ax[1].set_ylabel('X2')
ax[1].legend()
plt.show()

```



1.2 PART II

1.2.1 Q1

kNN classification (Attempt first for the Gaussian dataset, then repeat for the Laplacian dataset)

1. The k-Nearest Neighbors algorithm (kNN) assigns to a test point the most frequent label of its k closest examples in the training set.

(a) Write a function `kNNClassify` to generate predictions Y_p for the 2-class data

generated at Section 1. Pick a “reasonable” k .

(b) Evaluate the classification performance (prediction error) by comparing the predicted labels Y_p to the true labels Y_t

(c) Visualize the obtained results, e.g. by plotting the wrongly classified points using different colors/markers:

(d) Write a function to generate & visualize the decision regions of the 2D plane that are associated with each class, for a given classifier. Overlay the test points using scatter.

```
[ ]: n_neighbors = 15

def kNNClassify(X_train, Y_train, X_test, n_neighbors):

    distances = np.sqrt(np.sum((X_test[:, np.newaxis, :] - X_train[np.newaxis, :, :])**2, axis=2))

    kNN_loc = np.argsort(distances, axis=1)[:n_neighbors]

    k_NN_labels = Y_train[kNN_loc]

    Y_pred = np.array([np.argmax(np.bincount(x+1)) for x in k_NN_labels]) # +1
    ↪as bincount needs non-negative integers


    return Y_pred-1
```

```
[ ]: Y_pred = kNNClassify(X_train, Y_train, X_test, n_neighbors)
```

```
[ ]: wrong_loc = np.array([i for i in range(len(Y_pred)) if Y_pred[i] != Y_test[i]])

X_wrong = X_test[wrong_loc]
Y_wrong = Y_test[wrong_loc]

plt.scatter(X_wrong[Y_wrong == 1][:, 0], X_wrong[Y_wrong == 1][:, 1],
    ↪label='Class -1 pred 1 (FP)', marker='o')
plt.scatter(X_wrong[Y_wrong == -1][:, 0], X_wrong[Y_wrong == -1][:, 1],
    ↪label='Class 1 pred -1 (TN)', marker='x')
plt.legend()
plt.title('Wrongly classified samples')
plt.show()
```



```
[ ]: # Accuracy
accuracy = 100 * np.count_nonzero(Y_test == Y_pred)/len(Y_test)
print("Classification accuracy: ", accuracy)
```

Classification accuracy: 68.625

```
[ ]: def plot_decision_regions(X, Y, X_test, Y_test):
    h = 0.02

    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    Z = kNNClassify(X, Y, np.c_[xx.ravel(), yy.ravel()], n_neighbors)
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, cmap=plt.cm.summer, alpha=0.5)

    plt.scatter(X_test[Y_test == 1][:, 0], X_test[Y_test == 1][:, 1],
        label='Class +1', marker='o')
```

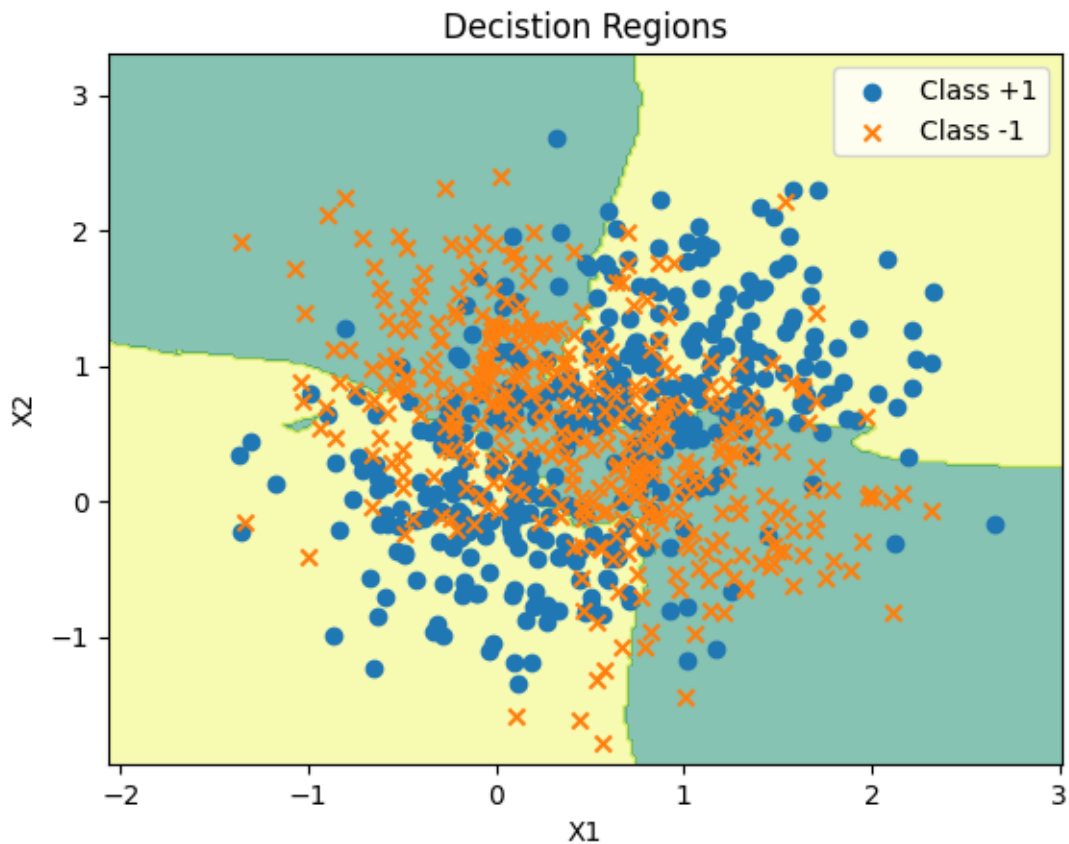
```

plt.scatter(X_test[Y_test == -1][:, 0], X_test[Y_test == -1][:, 1],
            label='Class -1', marker='x')

plt.xlabel("X1")
plt.ylabel("X2")
plt.legend()
plt.title("Decistion Regions")
plt.show()

plot_decision_regions(X_train, Y_train, X_test, Y_test)

```



1.2.2 Q2

- (a) Perform hold-out cross-validation by setting aside a fraction (ρ) of the training set for validation.

```

[ ]: hold_out_rho = 0.3
      k = 5
      n_rep = 10

```

```

def outCVkNN(X, Y, k, n_rep, hold_out_rho):
    n = len(X)
    n_hold_out = int(n*hold_out_rho)
    hold_out_accuracy = np.zeros(n_rep)
    train_accuracy = np.zeros(n_rep)
    for i in range(n_rep):
        shuffle_idx = np.random.permutation(n)
        X_train = X[shuffle_idx]
        Y_train = Y[shuffle_idx]
        X_hold_out = X_train[:n_hold_out]
        Y_hold_out = Y_train[:n_hold_out]
        X_train = X_train[n_hold_out:]
        Y_train = Y_train[n_hold_out:]
        Y_pred_hold_out = kNNClassify(X_train, Y_train, X_hold_out, k)
        hold_out_accuracy[i] = np.count_nonzero(Y_hold_out == Y_pred_hold_out)/
        ↪len(Y_hold_out)
        Y_pred_train = kNNClassify(X_train, Y_train, X_train, k)
        train_accuracy[i] = np.count_nonzero(Y_train == Y_pred_train)/
        ↪len(Y_train)

    return np.round([np.mean(hold_out_accuracy)*100, np.
    ↪mean(train_accuracy)*100], 2)

```

```

[ ]: aaa = []
    for k in range(1, 23, 2):
        knnout = outCVkNN(X_train, Y_train, k, n_rep, hold_out_rho)
        print("k = ", k, ": ", knnout)
        aaa.append(knnout)

```

```


k = 1 : [ 62.5 100. ]
k = 3 : [70.28 79.17]
k = 5 : [70.    77.62]
k = 7 : [71.39 79.29]
k = 9 : [71.67 79.52]
k = 11 : [73.33 78.81]
k = 13 : [76.11 75.71]
k = 15 : [71.11 77.02]
k = 17 : [71.39 75.36]
k = 19 : [69.44 74.76]
k = 21 : [72.22 73.69]

```

```

[ ]: plt.plot(range(1, 23, 2), [x[0] for x in aaa], label="Hold-out")
    plt.plot(range(1, 23, 2), [x[1] for x in aaa], label="Training")
    plt.legend()
    plt.show()

```



```
[ ]: k_best = np.argmax([x[0] for x in aaa])*2+1
      print('k =', k_best, 'is the best k for hold-out')
```

k = 13 is the best k for hold-out

- (b) How is the value of k affected by (percentage of points held out) and number of repetitions?
What does a large number of repetitions provide?

```
[ ]: # Rho vs k
best_k = []
for rho in [0.1, 0.3, 0.5, 0.7, 0.9]:
    # print("rho = ", rho, ": ", outCVkNN(X_train, Y_train, 5, 10, rho))
    aaa = []
    for k in range(1, 23, 2):
        knnout = outCVkNN(X_train, Y_train, k, n_rep, hold_out_rho)
        aaa.append(knnout)
    best_k.append(np.argmax([x[0] for x in aaa])*2+1)
    print('Rho = ' + str(rho) + ': k =', np.argmax([x[0] for x in aaa])*2+1,
          'is the best k for hold-out')
```


Rho = 0.1: k = 9 is the best k for hold-out

Rho = 0.3: k = 17 is the best k for hold-out

Rho = 0.5: k = 13 is the best k for hold-out

Rho = 0.7: k = 11 is the best k for hold-out
Rho = 0.9: k = 7 is the best k for hold-out

```
[ ]: plt.plot([0.1, 0.3, 0.5, 0.7, 0.9], best_k)
plt.xlabel("$\\rho$")
plt.ylabel("k")
plt.title("Best k vs $\\rho$")
plt.show()
```



```
[ ]: for n_rep in [1, 5, 10, 20, 30]:
      print("n_rep = ", n_rep, ": ", outCVkNN(X_train, Y_train, 5, n_rep, 0.3))
```

```
n_rep = 1 : [80.56 70.24]
n_rep = 5 : [75.   75.24]
n_rep = 10 : [70.   78.45]
n_rep = 20 : [69.44 79.05]
n_rep = 30 : [69.35 78.93]
```

```
[ ]: best_k_nrep = []
for n_rep in [1, 5, 10, 20, 30]:
```

```

# print("rho = ", rho, ": ", outCVkNN(X_train, Y_train, 5, 10, rho))
aaa = []
for k in range(1, 23, 2):
    knnout = outCVkNN(X_train, Y_train, k, n_rep, 0.3)
    aaa.append(knnout)
    best_k_nrep.append(np.argmax([x[0] for x in aaa])*2+1)
    # print('N_rep = ' + str(rho) + ': k = ', np.argmax([x[0] for x in
    ↪aaa])*2+1, 'is the best k for hold-out')
    print('For = ' + str(n_rep) + ' repetitions: k = ', np.argmax([x[0] for x
    ↪in aaa])*2+1, 'is the best')

```

```


For = 1 repetitions: k = 15 is the best
For = 5 repetitions: k = 11 is the best
For = 10 repetitions: k = 13 is the best
For = 20 repetitions: k = 11 is the best
For = 30 repetitions: k = 7 is the best

```

```

[ ]: #k vs n_rep
plt.plot([1, 5, 10, 20, 30], best_k)
plt.xlabel('n_rep')
plt.ylabel('k')
plt.title('Best k vs No.of repetitions')
plt.show()

```



(c) Apply the model obtained by cross-validation (i.e., best k) to the test set and check if there is an improvement on the classification error over the result of (1).

```
[ ]: acc_cv = 100 * np.count_nonzero(kNNClassify(X_train, Y_train, X_test, k_best)
    == Y_test)/len(Y_test)
# Accuracy
print("Classification accuracy: ", acc_cv)
```

Classification accuracy: 69.5

```
[ ]: if acc_cv > accuracy:
    print("Improved accuracy over (1)")
else:
    print("No improvement over (1)")
```

Improved accuracy over (1)

2 LAPLACE

2.1 PART I

(e) Repeat (a)-(d) for Laplace distribution.

```
[ ]: # TRAIN

train_samples_per_class = 30

# Set the variance for each class
variance = 0.3

n1a = np.random.laplace(loc=[0, 0], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
n1b = np.random.laplace(loc=[1, 1], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
n2a = np.random.laplace(loc=[0, 1], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))
n2b = np.random.laplace(loc=[1, 0], scale=np.sqrt(variance),
    ↪size=(train_samples_per_class, 2))

labels = np.array([1]*train_samples_per_class*2 +
    ↪[-1]*train_samples_per_class*2)

X_train = np.vstack([n1a, n1b, n2a, n2b])
Y_train = labels

shuffle_idx = np.random.permutation(X_train.shape[0])
X_train = X_train[shuffle_idx]
Y_train = Y_train[shuffle_idx]

# TEST

test_samples_per_class = 200

n1a = np.random.laplace(loc=[0, 0], scale=np.sqrt(variance),
    ↪size=(test_samples_per_class, 2))
n1b = np.random.laplace(loc=[1, 1], scale=np.sqrt(variance),
    ↪size=(test_samples_per_class, 2))
n2a = np.random.laplace(loc=[0, 1], scale=np.sqrt(variance),
    ↪size=(test_samples_per_class, 2))
n2b = np.random.laplace(loc=[1, 0], scale=np.sqrt(variance),
    ↪size=(test_samples_per_class, 2))
```

```

labels = np.array([1]*test_samples_per_class*2 + [-1]*test_samples_per_class*2)

X_test = np.vstack([n1a, n1b, n2a, n2b])
Y_test = labels

shuffle_idx = np.random.permutation(X_test.shape[0])
X_test = X_test[shuffle_idx]
Y_test = Y_test[shuffle_idx]


```

```

[ ]: fig, ax = plt.subplots(1, 2, figsize=(16, 6))
ax[0].scatter(X_train[Y_train == 1][:, 0], X_train[Y_train == 1][:, 1],
    ↪label='Class 1', marker='o')
ax[0].scatter(X_train[Y_train == -1][:, 0], X_train[Y_train == -1][:, 1],
    ↪label='Class -1', marker='x')
ax[0].set_title('Training set')
ax[0].set_xlabel('X1')
ax[0].set_ylabel('X2')
ax[0].legend()

ax[1].scatter(X_test[Y_test == 1][:, 0], X_test[Y_test == 1][:, 1],
    ↪label='Class 1', marker='o')
ax[1].scatter(X_test[Y_test == -1][:, 0], X_test[Y_test == -1][:, 1],
    ↪label='Class -1', marker='x')
ax[1].set_title('Test set')
ax[1].set_xlabel('X1')
ax[1].set_ylabel('X2')
ax[1].legend()
plt.show()

```



2.2 PART II

2.2.1 Q1

kNN classification

1. The k-Nearest Neighbors algorithm (kNN) assigns to a test point the most frequent label of its k closest examples in the training set.
 - (a) Write a function `kNNClassify` to generate predictions Y_p for the 2-class data generated at Section 1. Pick a “reasonable” k.
 - (b) Evaluate the classification performance (prediction error) by comparing the predicted labels Y_p to the true labels Y_t
 - (c) Visualize the obtained results, e.g. by plotting the wrongly classified points using different colors/markers:
 - (d) Write a function to generate & visualize the decision regions of the 2D plane that are associated with each class, for a given classifier. Overlay the test points using scatter.

```
[ ]: n_neighbors = 15

def kNNClassify(X_train, Y_train, X_test, n_neighbors):

    distances = np.sqrt(np.sum((X_test[:, np.newaxis, :] - X_train[np.newaxis, :
↪, :])**2, axis=2))

    kNN_loc = np.argsort(distances, axis=1)[: , :n_neighbors]

    k_NN_labels = Y_train[kNN_loc]

    Y_pred = np.array([np.argmax(np.bincount(x+1)) for x in k_NN_labels]) # +1↪
↪as bincount needs non-negative integers

    return Y_pred-1
```


```
[ ]: Y_pred = kNNClassify(X_train, Y_train, X_test, n_neighbors)
```

```
[ ]: wrong_loc = np.array([i for i in range(len(Y_pred)) if Y_pred[i] != Y_test[i]])

X_wrong = X_test[wrong_loc]
Y_wrong = Y_test[wrong_loc]

plt.scatter(X_wrong[Y_wrong == 1][:, 0], X_wrong[Y_wrong == 1][:, 1],↪
↪label='Class -1 pred 1 (FP)', marker='o')
plt.scatter(X_wrong[Y_wrong == -1][:, 0], X_wrong[Y_wrong == -1][:, 1],↪
↪label='Class 1 pred -1 (TN)', marker='x')
plt.legend()
```

```
plt.title('Wrongly classified samples')
plt.show()
```



```
[ ]: # Accuracy
accuracy = 100 * np.count_nonzero(Y_test == Y_pred)/len(Y_test)
print("Classification accuracy: ", accuracy)
```

Classification accuracy: 65.75

```
[ ]: def plot_decision_regions(X, Y, X_test, Y_test):
    h = 0.02

    x_min, x_max = X[:, 0].min() - 2, X[:, 0].max() + 2
    y_min, y_max = X[:, 1].min() - 2, X[:, 1].max() + 2
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    Z = kNNClassify(X, Y, np.c_[xx.ravel(), yy.ravel()], n_neighbors)
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, cmap=plt.cm.summer, alpha=0.5)
```


```

plt.scatter(X_test[Y_test == 1][:, 0], X_test[Y_test == 1][:, 1],
↪label='Class +1', marker='o')
plt.scatter(X_test[Y_test == -1][:, 0], X_test[Y_test == -1][:, 1],
↪label='Class -1', marker='x')

plt.xlabel("X1")
plt.ylabel("X2")
plt.legend()
plt.title("Decistion Regions")
plt.show()

plot_decision_regions(X_train, Y_train, X_test, Y_test)

```



2.2.2 Q2

- (a) Perform hold-out cross-validation by setting aside a fraction ($\frac{1}{n}$) of the training set for validation.


```
[ ]: hold_out_rho = 0.3
k = 5
n_rep = 10
def outCVkNN(X, Y, k, n_rep, hold_out_rho):
    n = len(X)
    n_hold_out = int(n*hold_out_rho)
    hold_out_accuracy = np.zeros(n_rep)
    train_accuracy = np.zeros(n_rep)
    for i in range(n_rep):
        shuffle_idx = np.random.permutation(n)
        X_train = X[shuffle_idx]
        Y_train = Y[shuffle_idx]
        X_hold_out = X_train[:n_hold_out]
        Y_hold_out = Y_train[:n_hold_out]
        X_train = X_train[n_hold_out:]
        Y_train = Y_train[n_hold_out:]
        Y_pred_hold_out = kNNClassify(X_train, Y_train, X_hold_out, k)
        hold_out_accuracy[i] = np.count_nonzero(Y_hold_out == Y_pred_hold_out)/
        ↪len(Y_hold_out)
        Y_pred_train = kNNClassify(X_train, Y_train, X_train, k)
        train_accuracy[i] = np.count_nonzero(Y_train == Y_pred_train)/
        ↪len(Y_train)


    return np.round([np.mean(hold_out_accuracy)*100, np.
    ↪mean(train_accuracy)*100], 2)
```

```
[ ]: aaa = []
for k in range(1, 23, 2):
    knnout = outCVkNN(X_train, Y_train, k, n_rep, hold_out_rho)
    print("k = ", k, ": ", knnout)
    aaa.append(knnout)
```

```
k = 1 : [ 67.5 100. ]
k = 3 : [67.78 82.38]
k = 5 : [74.17 80.6 ]
k = 7 : [75.28 78.1 ]
k = 9 : [75.  78.1]
k = 11 : [73.89 79.76]
k = 13 : [74.44 78.57]
k = 15 : [72.22 77.38]
k = 17 : [73.06 78.93]
k = 19 : [71.94 79.17]
k = 21 : [71.11 75.83]
```

```
[ ]: plt.plot(range(1, 23, 2), [x[0] for x in aaa], label="Hold-out")
plt.plot(range(1, 23, 2), [x[1] for x in aaa], label="Training")
```

```
plt.legend()
plt.show()
```



```
[ ]: k_best = np.argmax([x[0] for x in aaa])*2+1
print('k =', k_best, 'is the best k for hold-out')
```

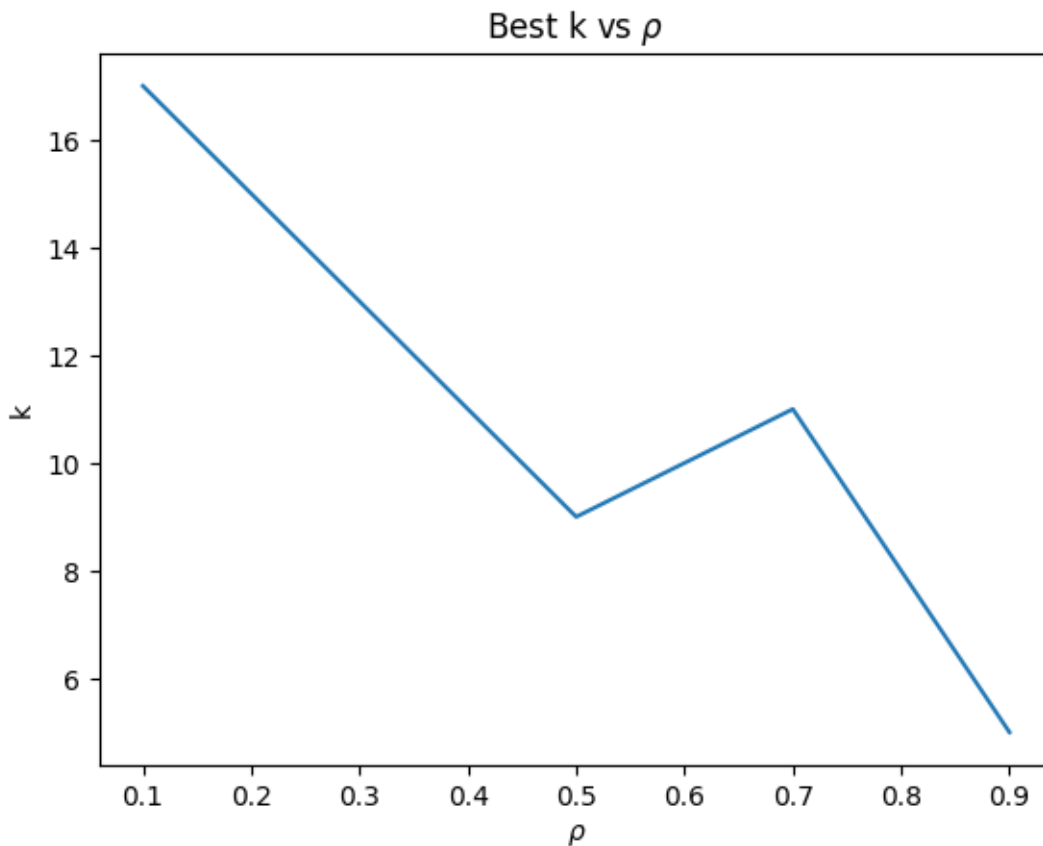
k = 7 is the best k for hold-out

- (b) How is the value of k affected by (percentage of points held out) and number of repetitions?
What does a large number of repetitions provide?

```
[ ]: # Rho vs k
best_k = []
for rho in [0.1, 0.3, 0.5, 0.7, 0.9]:
    # print("rho = ", rho, ": ", outCVkNN(X_train, Y_train, 5, 10, rho))
    aaa = []
    for k in range(1, 23, 2):
        knnout = outCVkNN(X_train, Y_train, k, n_rep, hold_out_rho)
        aaa.append(knnout)
    best_k.append(np.argmax([x[0] for x in aaa])*2+1)
    print('Rho = ' + str(rho) + ': k =', np.argmax([x[0] for x in aaa])*2+1,
    ↪ 'is the best k for hold-out')
```

Rho = 0.1: k = 17 is the best k for hold-out
 Rho = 0.3: k = 13 is the best k for hold-out
 Rho = 0.5: k = 9 is the best k for hold-out
 Rho = 0.7: k = 11 is the best k for hold-out
 Rho = 0.9: k = 5 is the best k for hold-out

```
[ ]: plt.plot([0.1, 0.3, 0.5, 0.7, 0.9], best_k)
plt.xlabel("$\\rho$")
plt.ylabel("k")
plt.title("Best k vs $\\rho$")
plt.show()
```



```
[ ]: for n_rep in [1, 5, 10, 20, 30]:
      print("n_rep = ", n_rep, ": ", outCVkNN(X_train, Y_train, 5, n_rep, 0.3))
```


```

n_rep = 1 : [77.78 79.76]
n_rep = 5 : [70.56 80.71]
n_rep = 10 : [70.28 81.55]
n_rep = 20 : [73.33 80. ]
n_rep = 30 : [74.17 80.04]
```

```
[ ]: best_k_nrep = []
for n_rep in [1, 5, 10, 20, 30]:
    # print("rho = ", rho, ": ", outCVkNN(X_train, Y_train, 5, 10, rho))
    aaa = []
    for k in range(1, 23, 2):
        knnout = outCVkNN(X_train, Y_train, k, n_rep, 0.3)
        aaa.append(knnout)
    best_k_nrep.append(np.argmax([x[0] for x in aaa])*2+1)
    # print('N_rep = ' + str(rho) + ': k = ', np.argmax([x[0] for x in
    ↪aaa])*2+1, 'is the best k for hold-out')
    print('For = ' + str(n_rep) + ' repetitions: k = ', np.argmax([x[0] for x
    ↪in aaa])*2+1, 'is the best')
```

```
For = 1 repetitions: k = 13 is the best
For = 5 repetitions: k = 7 is the best
For = 10 repetitions: k = 15 is the best
For = 20 repetitions: k = 11 is the best
For = 30 repetitions: k = 15 is the best
```

```
[ ]: #k vs n_rep
plt.plot([1, 5, 10, 20, 30], best_k)
plt.xlabel('n_rep')
plt.ylabel('k')
plt.title('Best k vs No.of repetitions')
plt.show()
```



(c) Apply the model obtained by cross-validation (i.e., best k) to the test set and check if there is an improvement on the classification error over the result of (1).

```
[ ]: acc_cv = 100 * np.count_nonzero(kNNClassify(X_train, Y_train, X_test, k_best)
    == Y_test)/len(Y_test)
# Accuracy
print("Classification accuracy: ", acc_cv)
```

Classification accuracy: 66.125

```
[ ]: if acc_cv > accuracy:
    print("Improved accuracy over (1)")
else:
    print("No improvement over (1)")
```

Improved accuracy over (1)